

```
In [4]: df = pd.read_csv("ratings.csv")
```

Out[4]:

	userid	movied	rating	timestamp
0	1	1	4.0	964982703
1	1	3	4.0	964981247
2	1	6	4.0	964982224
3	1	47	5.0	964983815
4	1	50	5.0	964982931
5	1	70	3.0	964982400
6	1	101	5.0	964980868
7	1	110	4.0	964982176
8	1	151	5.0	964980401
9	1	157	5.0	964984100
10	1	163	5.0	964983950
11	1	216	5.0	964981208
12	1	223	3.0	964980985
13	1	231	5.0	964981179
14	1	235	4.0	964980908
15	1	260	5.0	964981680
16	1	296	3.0	964982967
17	1	316	3.0	964982310
18	1	333	5.0	964981179
19	1	349	4.0	964982563
20	1	356	4.0	964980962
21	1	362	5.0	964982588
22	1	367	4.0	964981710
23	1	423	3.0	964982363
24	1	441	4.0	964980868
25	1	457	5.0	964981909
26	1	480	4.0	964982346
27	1	500	3.0	964981208
28	1	527	5.0	964980402
29	1	543	4.0	964981179
...				
100806	610	150401	3.0	1479543210
100807	610	152077	4.0	1493345517
100808	610	152081	4.0	1493485603
100809	610	152372	3.5	1493348841
100810	610	155064	3.5	1493348456
100811	610	156371	5.0	1479542831
100812	610	156726	4.0	1493348444
100813	610	157296	4.0	1493345563
100814	610	158238	5.0	1479545219
100815	610	158271	3.5	1479542491
100816	610	158732	3.5	1493348049
100817	610	158956	3.0	1493348947
100818	610	159093	3.0	1493347704
100819	610	160080	3.0	1493348031
100820	610	160341	2.5	1479545749
100821	610	160527	4.0	1479544998
100822	610	160571	3.0	1493348537
100823	610	160836	3.0	1493344794
100824	610	161582	4.0	1493347759
100825	610	161634	4.0	1493348362
100826	610	162350	3.5	1493348971
100827	610	163937	3.5	1493348748
100828	610	163981	3.5	1493305151
100829	610	164179	5.0	1493345631
100830	610	166528	4.0	1493379365
100831	610	166534	4.0	1493348402
100832	610	168248	5.0	1493385091
100833	610	168250	5.0	1494273047
100834	610	168252	5.0	1493463452
100835	610	170875	3.0	1493346451

100836 rows x 4 columns

## 1. Transforming Data

```
In [20]: userids = pd.unique(df['userId'])
movieids = pd.unique(df['movieId'])

matrix_T = pd.DataFrame(columns=movieids)
for user_id in userids:
    row = []
    temp_df = df[df['userId'] == user_id]
    for movie_id in movieids:
        if movie_id not in list(temp_df['movieId']):
            row.append(0)
        else:
            rate = temp_df[temp_df['movieId'] == movie_id].iloc[0,2]
            row.append(rate)
    matrix_T = matrix_T.append(pd.Series(row, index = movieids, ignore_index=True))

matrix_T.to_csv('matrix_T.csv', index = False)
print("Done.")

Done.
```

The user-movie ratings matrix is stored in "matrix\_T.csv".

```
In [2]: matrix_T = pd.read_csv("matrix_T.csv")
matrix_T
```

Out[2]:

[illegible]

4	4.0	0.0	0.0	0.0	4.0	0.0	0.0	4.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
5	0.0	5.0	4.0	4.0	1.0	0.0	0.0	5.0	4.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
6	4.5	0.0	0.0	0.0	4.5	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7	0.0	0.0	0.0	4.0	5.0	0.0	0.0	3.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
9	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
10	0.0	0.0	5.0	0.0	0.0	0.0	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
11	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
12	0.0	0.0	0.0	5.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
13	0.0	0.0	0.0	5.0	0.0	0.0	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
14	2.5	0.0	0.0	3.5	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
15	0.0	0.0	0.0	3.5	4.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
16	4.5	0.0	0.0	4.0	4.5	0.0	0.0	4.5	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
17	3.5	0.0	4.0	4.5	5.0	3.5	0.0	4.5	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
18	4.0	3.0	0.0	3.0	0.0	2.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
19	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
20	3.5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
21	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
22	0.0	0.0	4.0	0.0	4.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
23	0.0	0.0	4.5	0.0	4.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
24	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
25	0.0	0.0	0.0	4.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
26	3.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
27	0.0	0.0	3.5	3.0	3.5	0.0	0.0	3.5	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
28	0.0	0.0	0.0	0.0	3.5	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
29	0.0	0.0	0.0	0.0	0.0	0.0	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
580	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
581	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
582	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
583	5.0	0.0	0.0	5.0	0.0	0.0	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
584	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
585	0.0	0.0	0.0	0.0	0.0	0.0	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
586	5.0	0.0	0.0	0.0	5.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
587	0.0	3.0	5.0	3.0	5.0	0.0	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
588	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
589	4.0	3.0	3.5	3.0	4.5	0.0	0.0	4.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
590	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
591	0.0	0.0	3.0	0.0	0.0	0.0	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
592	0.0	0.0	0.0	0.0	4.5	0.0	0.0	3.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
593	0.0	4.0	0.0	0.0	0.0	3.5	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
594	0.0	0.0	0.0	0.0	5.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
595	4.0	0.0	0.0	0.0	3.5	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
596	4.0	0.0	3.0	4.0	5.0	2.0	5.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
597	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
598	3.0	1.5	4.5	4.0	3.5	3.5	2.5	3.5	0.0	1.5	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
599	2.5	0.0	0.0	4.0	0.0	0.0	4.5	2.0	3.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
600	4.0	0.0	0.0	4.0	5.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
601	0.0	0.0	3.0	5.0	5.0	0.0	0.0	5.0	4.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
602	4.0	0.0	4.0	0.0	0.0	4.0	4.0	1.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
603	3.0	0.0	3.0	0.0	0.0	0.0	0.0	3.0	3.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
604	4.0	0.0	0.0	0.0	0.0	0.0	0.0	3.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
605	2.5	0.0	0.0	3.0	4.5	4.0	0.0	3.5	0.0	4.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
606	4.0	0.0	0.0	0.0	0.0	0.0	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
607	2.5	2.0	0.0	4.5	4.5	3.0	0.0	4.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
608	3.0	0.0	0.0	0.0	0.0	0.0	0.0	3.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
609	5.0	0.0	5.0	5.0	4.0	4.0	0.0	4.5	0.0	0.0	...	3.0	3.5	3.5	3.5	3.5	2.5	4.5	3.0
610	0.0	0.0	0.0	0.0	0.0	0.0	0.0	5.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

610 rows x 9724 columns

```
In [15]: matrix_T.shape
```

```
Out[15]: (610, 9724)
```

Since every user has only rated a very small part of the movie and the missing values are replaced by 0 in the user-movie ratings matrix, the most common entry in this constructed matrix is 0. We call a matrix like this a recommender system.

## 2. Principle Component Analysis

(a)

```
In [4]: # Transform the data
matrix = matrix_T.T
```


```
In [33]: # Mean center the data
m_scaled = preprocessing.scale(matrix, with_std=False)
```

(b)

```
In [34]: # Apply PCA with number of components k = 2
pca = decomposition.PCA(n_components=2)
pca.fit(m_scaled)
m_trans = pca.transform(m_scaled)
print(m_trans.shape)

# Plot the new representations of the movies with a scatter plot.
plt.scatter(m_trans[:,0],m_trans[:,1])
plt.show()
```

(9724, 2)



(c)

```
In [35]: pca.explained_variance_ratio_
Out[35]: array([0.17620942, 0.04189505])
```

From the plot, we can see clearly that data spare more widely on the x-axis, hile the explained variance ratio indicated that the first dimension has greater variance.

(d)

```
In [31]: variance_sum = 0
        k=1
        while variance_sum < 0.8:
            pca = decomposition.PCA(n_components=k)
            pca.fit(X_scaled)
            variance_sum = sum(pca.explained_variance_ratio_)
            k += 1
        print("k = " +str(k))
        print("Total variance: " + str(variance_sum))

k = 155
Total variance: 0.8003133112551502
```

With k=2, only 21.7% of data is explained. To explain 80% of the variance of the data, 155 principle components are needed. When it comes to visualization, we notice that it's impossible to plot the data with 155 dimensions, while k=2 can be easily plotted. Therefore, we can see that when k gets larger, it can explain more data, but it might not be able to visualize.

### 3. K-Mean

(a)


```
In [10]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.cluster import KMeans
import matplotlib

ks=[2, 4, 8, 16, 32]
inertias = []
for k in ks:
    kmeans = KMeans(n_clusters=k)
    kmeans.fit(matrix)
    y_kmeans = kmeans.predict(matrix)
    inertia = kmeans.inertia_
    inertias.append([k, inertia])
    print(inertia)

inertias = np.array(inertias)
plt.scatter(inertias[:, 0], inertias[:, 1])
plt.plot(inertias[:, 0], inertias[:, 1])

1124737.0197068886
1052374.2609276343
993351.0115743998
942772.9606929905
898671.0869614722

Out[10]: [ <matplotlib.lines.Line2D at 0x1635d452b00>]
```



Add more values of k


```
In [11]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.cluster import KMeans
import matplotlib

ks=[2, 4, 8, 16, 24, 32, 48, 64, 96, 128]
inertias = []
for k in ks:
    kmeans = KMeans(n_clusters=k)
    kmeans.fit(matrix)
    y_kmeans = kmeans.predict(matrix)
    inertia = kmeans.inertia_
    inertias.append([k, inertia])
    print(inertia)

inertias = np.array(inertias)
plt.scatter(inertias[:, 0], inertias[:, 1])
plt.plot(inertias[:, 0], inertias[:, 1])

1124737.0197068886
1052373.6109772539
964368.0794482677
943864.4636031041
938651.1631696054
887832.3477151245
855284.4860510383
823387.3922342558
777122.7756901473
746326.0422153994

Out[11]: [ <matplotlib.lines.Line2D at 0x1635d45ada0>]
```



(b)

Choose k = 24.  
From the plot, we can see that 24 is the ideal number of k. It is at the "elbow" of the plot. Therefore, it provides a good tradeoff between accuracy (low value of inertia) and complexity (small value of k).

(c)

Movies could be clustered together by the type of movies, i.e. action, sci-fi, horror...  
This type of interpretation might not always be possible. Since by choosing different k, the number of clusters is different, which can merge several types of movies into a single type of movie with a small k. For example, if k is really small, comedy, romance, and animation could be grouped together as "fun" movie. While, horror, war, and crime may be clustered as "sad" movie. However, if k a large number, one type of movie might be refined into several types of movies. For

example, action moves can be interpreted into different types of rs poisons or actors.


## 4. Singular Value Decomposition

(a)

```
In [19]: from sklearn.decomposition import TruncatedSVD
from sklearn.random_projection import sparse_random_matrix

svd = TruncatedSVD(n_components=32)
svd.fit(matrix)
singular = svd.singular_values_
plt.plot(singular)
```

Out[19]: [matplotlib.lines.Line2D at 0x1635c7460480]



(b)

```
In [14]: variances = []
for k in [2, 4, 8, 16, 32]:
    svd = TruncatedSVD(n_components=k)
    svd.fit(matrix)
    variance = sum(svd.explained_variance_ratio_)
    variances.append((k, variance))
    print(str(k) + ":")
    print(svd.singular_values_)
    print(sum(svd.explained_variance_ratio_))
    print("\n")

variances = np.array(variances)
plt.scatter(variances[:, 0], variances[:, 1])
plt.plot(variances[:, 0], variances[:, 1])

2:
[534.41989777 231.23661061]
0.21642917529752204


4:
[534.41989777 231.236611 191.15084454 170.42188613]
0.2841428516880421

8:
[534.41989777 231.23661123 191.15084979 170.42239983 154.5522061]
147.33464029 135.64513529 122.61186255]
0.32499747809948353

16:
[534.41989777 231.2366114 191.1508759 170.42249307 154.55274805]
147.33552678 135.65525637 122.66035184 121.43769458 113.08261962]
109.59533707 107.83483322 105.85954835 101.77109818 99.57623689]
99.06365971]
0.39724086776976014

32:
[534.41989777 231.23661141 191.15087615 170.42250816 154.55294482]
147.33574096 135.6555122 122.66253061 121.44176937 113.11034896]
109.5964902 107.92580094 105.96758061 102.04653537 99.86115755]
99.27106335 97.10015173 93.29345448 92.24387347 90.89076604]
90.15260453 88.71268596 87.06893178 85.85554522 84.90663537]
82.52429351 82.03964364 81.32257963 80.60800026 79.10004555]
78.48117571 76.88147308]
0.4896308009255997
```

Out[14]: [matplotlib.lines.Line2D at 0x1635c5d8da00]




When increasing k in k-mean, the inertia values decreases which means that entities are getting closer to its cluster's center point. As for SVD, increasing k can increase explained variance ratio. In that case, more data can be explained after dimensionality reduction. I chose k = 24 in the previous question. From the plot above, we can see that around 45% of data can be explained by this number of components.

(c)

```
In [21]: svd = TruncatedSVD(n_components=2)
svd.fit(matrix)
a_reduced = svd.fit_transform(matrix)
a_reduced
plt.scatter(a_reduced[:, 0], a_reduced[:, 1])


Out[21]: [matplotlib.collections.PathCollection at 0x1635c7d35f80]
```



(d)

```
In [22]: kmeans = KMeans(n_clusters=24)
kmeans.fit(matrix)
y_kmeans = kmeans.predict(matrix)
plt.scatter(a_reduced[:, 0], a_reduced[:, 1], c=y_kmeans, cmap=plt.cm.RdYlGn)

Out[22]: [matplotlib.collections.PathCollection at 0x1635c5ac3470]
```



From the plot above, we can see that data transformed by SVD has almost the same shape as data transformed by PCA when the number of components is 2. Besides, movies are colored by the cluster memberships discovered by k-mean. It obvious that movies in the same cluster are grouped together. This support that clustering works well.